

```
pip install pandas-profiling
```

```
Looking in indexes: https://pypi.org/simple, https://us-  
python.pkg.dev/colab-wheels/public/simple/
```

```
Collecting pandas-profiling
```

```
  Downloading pandas_profiling-3.6.6-py2.py3-none-any.whl (324 kB)
```

```
----- 324.4/324.4 kB 11.5 MB/s eta  
0:00:00
```

```
----- 345.9/345.9 kB 36.5 MB/s eta
```

```
0:00:00
```

```
anylinux_2_17_x86_64.manylinux2014_x86_64.whl (679 kB)
```

```
----- 679.8/679.8 kB 52.6 MB/s eta
```

```
0:00:00
```

```
Requirement already satisfied: seaborn<0.13,>=0.10.1 in
```

```
/usr/local/lib/python3.9/dist-packages (from ydata-profiling->pandas-  
profiling) (0.12.2)
```

```
Requirement already satisfied: pandas!=1.4.0,<1.6,>1.1 in
```

```
/usr/local/lib/python3.9/dist-packages (from ydata-profiling->pandas-  
profiling) (1.5.3)
```

```
Collecting tqdm<4.65,>=4.48.2
```

```
  Downloading tqdm-4.64.1-py2.py3-none-any.whl (78 kB)
```

```
----- 78.5/78.5 kB 9.7 MB/s eta
```

```
0:00:00
```

```
Requirement already satisfied: pydantic<1.11,>=1.8.1 in
```

```
/usr/local/lib/python3.9/dist-packages (from ydata-profiling->pandas-  
profiling) (1.10.7)
```

```
Collecting imagehash==4.3.1
```

```
  Downloading ImageHash-4.3.1-py2.py3-none-any.whl (296 kB)
```

```
----- 296.5/296.5 kB 36.3 MB/s eta
```

```
0:00:00
```

```
Requirement already satisfied: numpy<1.24,>=1.16.0 in
```

```
/usr/local/lib/python3.9/dist-packages (from ydata-profiling->pandas-  
profiling) (1.22.4)
```

```
Requirement already satisfied: statsmodels<0.14,>=0.13.2 in
```

```
/usr/local/lib/python3.9/dist-packages (from ydata-profiling->pandas-  
profiling) (0.13.5)
```

```
Requirement already satisfied: requests<2.29,>=2.24.0 in
```

```
/usr/local/lib/python3.9/dist-packages (from ydata-profiling->pandas-  
profiling) (2.27.1)
```

```
Collecting typeguard<2.14,>=2.13.2
```

```
  Downloading typeguard-2.13.3-py3-none-any.whl (17 kB)
```

```
Requirement already satisfied: jinja2<3.2,>=2.11.1 in
```

```
/usr/local/lib/python3.9/dist-packages (from ydata-profiling->pandas-  
profiling) (3.1.2)
```

```
Collecting scipy<1.10,>=1.4.1
```

```
  Downloading scipy-1.9.3-cp39-cp39-
```

```
anylinux_2_17_x86_64.manylinux2014_x86_64.whl (33.8 MB)
```

```
----- 33.8/33.8 MB 19.5 MB/s eta
```

```
0:00:00
```

```
age_path]==0.7.5
```

```
Downloading visions-0.7.5-py3-none-any.whl (102 kB)
102.7/102.7 kB 12.4 MB/s eta
0:00:00
atplotlib<3.7,>=3.2
  Downloading matplotlib-3.6.3-cp39-cp39-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl (11.8 MB)
11.8/11.8 MB 86.6 MB/s eta
0:00:00
ultimethod<1.10,>=1.4
  Downloading multimethod-1.9.1-py3-none-any.whl (10 kB)
Collecting htmlmin==0.1.12
  Downloading htmlmin-0.1.12.tar.gz (19 kB)
  Preparing metadata (setup.py) ... ent already satisfied:
PyYAML<6.1,>=5.0.0 in /usr/local/lib/python3.9/dist-packages (from
ydata-profiling->pandas-profiling) (6.0)
Requirement already satisfied: PyWavelets in
/usr/local/lib/python3.9/dist-packages (from imagehash==4.3.1->ydata-
profiling->pandas-profiling) (1.4.1)
Requirement already satisfied: pillow in
/usr/local/lib/python3.9/dist-packages (from imagehash==4.3.1->ydata-
profiling->pandas-profiling) (8.4.0)
Requirement already satisfied: networkx>=2.4 in
/usr/local/lib/python3.9/dist-packages (from
visions[type_image_path]==0.7.5->ydata-profiling->pandas-profiling)
(3.1)
Collecting tangled-up-in-unicode>=0.0.4
  Downloading tangled_up_in_unicode-0.2.0-py3-none-any.whl (4.7 MB)
4.7/4.7 MB 73.3 MB/s eta
0:00:00
ent already satisfied: attrs>=19.3.0 in /usr/local/lib/python3.9/dist-
packages (from visions[type_image_path]==0.7.5->ydata-profiling-
>pandas-profiling) (23.1.0)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.9/dist-packages (from jinja2<3.2,>=2.11.1-
>ydata-profiling->pandas-profiling) (2.1.2)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.9/dist-packages (from matplotlib<3.7,>=3.2-
>ydata-profiling->pandas-profiling) (2.8.2)
Requirement already satisfied: pyparsing>=2.2.1 in
/usr/local/lib/python3.9/dist-packages (from matplotlib<3.7,>=3.2-
>ydata-profiling->pandas-profiling) (3.0.9)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.9/dist-packages (from matplotlib<3.7,>=3.2-
>ydata-profiling->pandas-profiling) (4.39.3)
Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.9/dist-packages (from matplotlib<3.7,>=3.2-
>ydata-profiling->pandas-profiling) (0.11.0)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.9/dist-packages (from matplotlib<3.7,>=3.2-
>ydata-profiling->pandas-profiling) (23.1)
```

Requirement already satisfied: contourpy>=1.0.1 in
 /usr/local/lib/python3.9/dist-packages (from matplotlib<3.7,>=3.2-
 >ydata-profiling->pandas-profiling) (1.0.7)

Requirement already satisfied: kiwisolver>=1.0.1 in
 /usr/local/lib/python3.9/dist-packages (from matplotlib<3.7,>=3.2-
 >ydata-profiling->pandas-profiling) (1.4.4)

Requirement already satisfied: pytz>=2020.1 in
 /usr/local/lib/python3.9/dist-packages (from pandas!=1.4.0,<1.6,>1.1-
 >ydata-profiling->pandas-profiling) (2022.7.1)

Requirement already satisfied: joblib>=0.14.1 in
 /usr/local/lib/python3.9/dist-packages (from phik<0.13,>=0.11.1-
 >ydata-profiling->pandas-profiling) (1.2.0)

Requirement already satisfied: typing-extensions>=4.2.0 in
 /usr/local/lib/python3.9/dist-packages (from pydantic<1.11,>=1.8.1-
 >ydata-profiling->pandas-profiling) (4.5.0)

Requirement already satisfied: idna<4,>=2.5 in
 /usr/local/lib/python3.9/dist-packages (from requests<2.29,>=2.24.0-
 >ydata-profiling->pandas-profiling) (3.4)

Requirement already satisfied: urllib3<1.27,>=1.21.1 in
 /usr/local/lib/python3.9/dist-packages (from requests<2.29,>=2.24.0-
 >ydata-profiling->pandas-profiling) (1.26.15)

Requirement already satisfied: certifi>=2017.4.17 in
 /usr/local/lib/python3.9/dist-packages (from requests<2.29,>=2.24.0-
 >ydata-profiling->pandas-profiling) (2022.12.7)

Requirement already satisfied: charset-normalizer~=2.0.0 in
 /usr/local/lib/python3.9/dist-packages (from requests<2.29,>=2.24.0-
 >ydata-profiling->pandas-profiling) (2.0.12)

Requirement already satisfied: patsy>=0.5.2 in
 /usr/local/lib/python3.9/dist-packages (from
 statsmodels<0.14,>=0.13.2->ydata-profiling->pandas-profiling) (0.5.3)

Requirement already satisfied: six in /usr/local/lib/python3.9/dist-
 packages (from patsy>=0.5.2->statsmodels<0.14,>=0.13.2->ydata-
 profiling->pandas-profiling) (1.16.0)

Building wheels for collected packages: htmlmin
 Building wheel for htmlmin (setup.py) ... lmin: filename=htmlmin-
 0.1.12-py3-none-any.whl size=27096
 sha256=a94b02692dc5b1b0e3514871d9628470911ea1801a8a7fe66699ba0c2b5ea31
 3

Stored in directory:
 /root/.cache/pip/wheels/1d/05/04/c6d7d3b66539d9e659ac6dfe81e2d0fd4c1a8
 316cc5a403300

Successfully built htmlmin

Installing collected packages: htmlmin, typeguard, tqdm, tangled-up-
 in-unicode, scipy, multimethod, matplotlib, imagehash, visions, phik,
 ydata-profiling, pandas-profiling

Attempting uninstall: tqdm
 Found existing installation: tqdm 4.65.0
 Uninstalling tqdm-4.65.0:
 Successfully uninstalled tqdm-4.65.0

Attempting uninstall: scipy

```
Found existing installation: scipy 1.10.1
Uninstalling scipy-1.10.1:
  Successfully uninstalled scipy-1.10.1
Attempting uninstall: matplotlib
Found existing installation: matplotlib 3.7.1
Uninstalling matplotlib-3.7.1:
  Successfully uninstalled matplotlib-3.7.1
Successfully installed htmlmin-0.1.12 imagehash-4.3.1 matplotlib-3.6.3
multimethod-1.9.1 pandas-profiling-3.6.6 phik-0.12.3 scipy-1.9.3
tangled-up-in-unicode-0.2.0 tqdm-4.64.1 typeguard-2.13.3 visions-0.7.5
ydata-profiling-4.1.2
```

```
{"pip_warning":{"packages":["matplotlib","mpl_toolkits"]}}
```

```
import pandas_profiling
```

```
<ipython-input-2-6a00893fb3e1>:1: DeprecationWarning: `import
pandas_profiling` is going to be deprecated by April 1st. Please use
`import ydata_profiling` instead.
import pandas_profiling
```

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.linear_model import Ridge, Lasso, RidgeCV, LassoCV,
ElasticNet, ElasticNetCV, LogisticRegression
```

```
from sklearn.model_selection import train_test_split
```

```
from statsmodels.stats.outliers_influence import
variance_inflation_factor
```

```
from sklearn.metrics import accuracy_score, confusion_matrix,
roc_curve, roc_auc_score
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
import pickle
```

```
df =
pd.read_csv("https://raw.githubusercontent.com/plotly/datasets/master/
diabetes.csv")
```

```
df
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
\						
0	6	148	72	35	0	33.6

1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1
..
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
from pandas_profiling import ProfileReport
```

```
ProfileReport(df)
```

```
{"model_id": "0568067ef96f434ea205af8ee9005eae", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "30d9d8d1a94846188895959eed53fb06", "version_major": 2, "version_minor": 0}
```


```
{"model_id": "29eb2d9770de4d3bb97a247c09348dda", "version_major": 2, "version_minor": 0}
```

<IPython.core.display.HTML object>

```
df['BMI'] = df['BMI'].replace(0, df['BMI'].mean())
df.columns
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
       'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
df['BloodPressure'] = df['BloodPressure'].replace(0,
df['BloodPressure'].mean())
df['Glucose'] = df['Glucose'].replace(0, df['Glucose'].mean())
df['Insulin'] = df['Insulin'].replace(0, df['Insulin'].mean())
df['SkinThickness'] = df['SkinThickness'].replace(0,
df['SkinThickness'].mean())

Box Plot to show outliers
fig, ax = plt.subplots(figsize = (20, 20))
sns.boxplot(data = df, ax = ax)


<Axes: >
```



```
q = df['Insulin'].quantile(.70)
df_new = df[df['Insulin'] < q]

fig, ax = plt.subplots(figsize = (20, 20))
sns.boxplot(data = df_new, ax = ax)

<Axes: >
```



```
q = df['Pregnancies'].quantile(.98)
df_new = df[df['Pregnancies'] < q]
```

```
q = df['BMI'].quantile(.99)
df_new = df_new[df_new['BMI'] < q]
```

```
q = df['SkinThickness'].quantile(.99)
df_new = df_new[df_new['SkinThickness'] < q]
```

```
q = df['Insulin'].quantile(.99)
df_new = df_new[df_new['Insulin'] < q]
```

```
q = df['Age'].quantile(.99)
df_new = df_new[df_new['Age'] < q]
```

```
df_new
```



BMI \	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
0	6	148.0	72.0	35.000000	79.799479
33.6					
1	1	85.0	66.0	29.000000	79.799479
26.6					
2	8	183.0	64.0	20.536458	79.799479
23.3					
3	1	89.0	66.0	23.000000	94.000000
28.1					
4	0	137.0	40.0	35.000000	168.000000
43.1					
..
...					
763	10	101.0	76.0	48.000000	180.000000
32.9					
764	2	122.0	70.0	27.000000	79.799479
36.8					
765	5	121.0	72.0	23.000000	112.000000
26.2					
766	1	126.0	60.0	20.536458	79.799479
30.1					
767	1	93.0	70.0	31.000000	79.799479
30.4					

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[716 rows x 9 columns]

```
fig, ax = plt.subplots(figsize = (20, 20))
sns.boxplot(data = df_new, ax = ax)
```

<Axes: >



ProfileReport(df_new)

```
{"model_id":"136761633ad04989af90297d87c5c8a2","version_major":2,"version_minor":0}
```

```
{"model_id":"a0bd6f0f2002480eab7aba9f415b3d32","version_major":2,"version_minor":0}
```

```
{"model_id":"b3c5b957590342aabb18be94efd07a29","version_major":2,"version_minor":0}
```

<IPython.core.display.HTML object>

df_new

```

Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin
BMI \

```

0	6	148.0	72.0	35.000000	79.799479
33.6					
1	1	85.0	66.0	29.000000	79.799479
26.6					
2	8	183.0	64.0	20.536458	79.799479
23.3					
3	1	89.0	66.0	23.000000	94.000000
28.1					
4	0	137.0	40.0	35.000000	168.000000
43.1					
..
...					
763	10	101.0	76.0	48.000000	180.000000
32.9					
764	2	122.0	70.0	27.000000	79.799479
36.8					
765	5	121.0	72.0	23.000000	112.000000
26.2					
766	1	126.0	60.0	20.536458	79.799479
30.1					
767	1	93.0	70.0	31.000000	79.799479
30.4					

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[716 rows x 9 columns]

y = df_new['Outcome']

y

0	1
1	0
2	1
3	0
4	1
..	
763	0
764	0
765	0

```

766     1
767     0
Name: Outcome, Length: 716, dtype: int64

```

```

x = df_new.drop(columns=['Outcome'])

```

```

x

```

BMI \	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
0	6	148.0	72.0	35.000000	79.799479
33.6					
1	1	85.0	66.0	29.000000	79.799479
26.6					
2	8	183.0	64.0	20.536458	79.799479
23.3					
3	1	89.0	66.0	23.000000	94.000000
28.1					
4	0	137.0	40.0	35.000000	168.000000
43.1					
..
...					
763	10	101.0	76.0	48.000000	180.000000
32.9					
764	2	122.0	70.0	27.000000	79.799479
36.8					
765	5	121.0	72.0	23.000000	112.000000
26.2					
766	1	126.0	60.0	20.536458	79.799479
30.1					
767	1	93.0	70.0	31.000000	79.799479
30.4					

	DiabetesPedigreeFunction	Age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33
..
763	0.171	63
764	0.340	27
765	0.245	30
766	0.349	47
767	0.315	23

```

[716 rows x 8 columns]

```


```

scaler = StandardScaler()
ProfileReport(pd.DataFrame(scaler.fit_transform(x)))
x_scaled = scaler.fit_transform(x)

```

```
df_new_scaler = pd.DataFrame(scaler.fit_transform(df_new))
fig, ax = plt.subplots(figsize = (20, 20))
sns.boxplot(data = df_new_scaler, ax = ax)
```

<Axes: >



Dataset is dispersed equally

x_scaled

```
array([[ 0.80822764,  0.9059344 ,  0.003933  , ...,  0.26388706,
         0.5060621 ,  1.60917704],
       [-0.86043151, -1.19643634, -0.49855498, ..., -0.67766768,
        -0.35435232, -0.12615744],
       [ 1.4756913 ,  2.07391814, -0.66605097, ..., -1.12154349,
         0.64634705, -0.03482405],
       ...,
       [ 0.47449581,  0.00491836,  0.003933  , ..., -0.73147081,
```

```

        -0.68480133, -0.21749083],
        [-0.86043151,  0.17177319, -1.00104295, ..., -0.20689031,
        -0.3605872 ,  1.33517686],
        [-0.86043151, -0.92946863, -0.16356299, ..., -0.16653796,
        -0.46658028, -0.85682459]])

```

y

```

0      1
1      0
2      1
3      0
4      1
..
763    0
764    0
765    0
766    1
767    0

```

Name: Outcome, Length: 716, dtype: int64

Checking Colinearity

```

def vif_score(x):
    scaler = StandardScaler()
    arr = scaler.fit_transform(x)
    return pd.DataFrame([x.columns[i],
    variance_inflation_factor(arr,i)] for i in range(arr.shape[1])),
    columns = ["FEATURE", "VIF_SCORE"])

```

vif_score(x)

	FEATURE	VIF_SCORE
0	Pregnancies	1.474617
1	Glucose	1.363812
2	BloodPressure	1.226262
3	SkinThickness	1.420421
4	Insulin	1.305182
5	BMI	1.438581
6	DiabetesPedigreeFunction	1.041569
7	Age	1.681447

```

x_train, x_test, y_train, y_test = train_test_split(x_scaled, y,
test_size = .20, random_state = 144)

```

x_train

```

array([[ -0.19296785,  1.64009561, -0.66605097, ...,  0.3849441 ,
        -0.3387651 , -0.21749083],
        [-1.19416334,  1.33975693, -1.83852292, ..., -1.30985444,
        -0.65674434,  2.97917794],
        [ 1.14195947,  2.44099875, -0.33105899, ...,  0.57325505,

```

```

0.87392043, 0.7871765 ],
...,
[-0.86043151, -1.39666213, 0.67391697, ..., -0.83907707,
 0.36889458, -0.94815798],
[ 0.14076398, 0.37199897, 1.17640494, ..., -0.48935673,
 -0.14236616, 2.79651116],
[ 0.47449581, 0.77245054, 0.84141296, ..., 0.04867455,
 -0.03949052, 2.33984419]])

```

x_test

```

array([[ 0.80822764,  0.77245054,  0.003933 , ...,  0.30423941,
        -0.65362689,  0.6958431 ],
       [ 0.47449581, -1.09632345, -0.49855498, ..., -0.97358489,
        -0.38240931, -0.21749083],
       [-0.86043151, -1.22980731, -0.66605097, ...,  0.70776287,
        0.01974091, -0.40015762],
       ...,
       [-1.19416334, -0.32879128, -0.58230297, ..., -0.94668332,
        0.60893773, -0.12615744],
       [-1.19416334, -0.09519453,  1.00890895, ...,  1.90488247,
        0.26913639, -0.12615744],
       [ 0.47449581, -0.16193646,  0.17142899, ...,  0.08902689,
        0.60893773,  0.23917614]])

```

y_train

```

132    1
294    0
185    1
643    0
634    0
..
693    1
456    0
232    0
479    0
388    1
Name: Outcome, Length: 572, dtype: int64

```

y_test

```

95     0
52     0
554    0
473    0
690    0
..
214    1
474    0
529    0

```

```

16      1
386     1
Name: Outcome, Length: 144, dtype: int64

x_test[0]

array([ 0.80822764,  0.77245054,  0.003933 ,  0.10259685,
        1.48329699,
        0.30423941, -0.65362689,  0.6958431 ])

logr_liblinear = LogisticRegression(verbose = 1, solver='liblinear')
logr.fit(x_train, y_train)

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1
concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.0s finished

LogisticRegression(verbose=1)

logr.predict([x_test[1]])

array([0])

logr.predict_proba([x_test[1]])

array([[0.92571637, 0.07428363]])

logr.predict_log_proba([x_test[1]])

array([[ -0.07718739, -2.59986461]])

type(y_test)

pandas.core.series.Series

y_test.iloc[0]

0

y_test

95      0
52      0
554     0
473     0
690     0
..
214     1
474     0
529     0
16      1
386     1
Name: Outcome, Length: 144, dtype: int64

```



```

logr = LogisticRegression(verbose = 1)
logr.fit(x_train, y_train)
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1
concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.0s finished
LogisticRegression(verbose=1)
logr
LogisticRegression(verbose=1)
logr_liblinear
LogisticRegression(solver='liblinear', verbose=1)
logr_liblinear.fit(x_train, y_train)
[LibLinear]
LogisticRegression(solver='liblinear', verbose=1)
y_pred_liblinear = logr_liblinear.predict(x_test)
y_pred_liblinear
array([1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0,
1,
      1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0,
0,
      0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0,
1,
      0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1,
0,
      1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0,
1,
      1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0,
0,
      0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0])
y_pred_default = logr.predict(x_test)
y_pred_default
array([1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0,
1,
      1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0,
0,
      0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0,
1,
      0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1,
0,

```

```

1,      1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0,
1,      1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0,
0,      0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0])
confusion_matrix(y_test, y_pred_liblinear)
array([[87,  7],
       [19, 31]])
confusion_matrix(y_test, y_pred_default)
array([[87,  7],
       [19, 31]])
auc = roc_auc_score(y_test, y_pred_liblinear)
roc_auc_score(y_test, y_pred_default)
0.7727659574468084

```

```

fpr, tpr, thresholds = roc_curve(y_test, y_pred_liblinear)
roc_curve(y_test, y_pred_liblinear)
(array([0.          , 0.07446809, 1.          ]),
 array([0.        , 0.62, 1.        ]),
 array([2, 1, 0]))

plt.plot(fpr, tpr, color = 'orange', label='ROC')
plt.plot([0, 1], [0,1], color='darkblue', linestyle='--', label='ROC
curve (area = %0.2f)' % auc)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Reciever Operating Characteristic (ROC) Curve')
plt.legend()
plt.show()

```

